

Worker goal reward (cosine_max) across training, mean $\pm$ std over seeds

